

# Congestion Control III

Performance improvements: fast retransmit and fast recovery


# Three Mechanisms

- Fast retransmit (TCP Tahoe): don't wait for a timeout to retransmit a missing segment if you receive a triple duplicate acknowledgement
  - ▶ Only drop back to slow start state on a timeout
- Fast recovery (TCP Reno): halve the congestion window (don't set it to 1) on triple duplicate acknowledgements
- Fast recovery (TCP Reno): while in fast recovery state, inflate the congestion window as acknowledgements arrive, to keep data flowing
  - ▶ Each duplicate ack increases congestion window by 1
  - ▶ If the old window is  $c$ , then the new window is  $c/2$
  - ▶ Receiving  $c$  acks will increase window size to  $3c/2$  -- can send  $c/2$  new segments

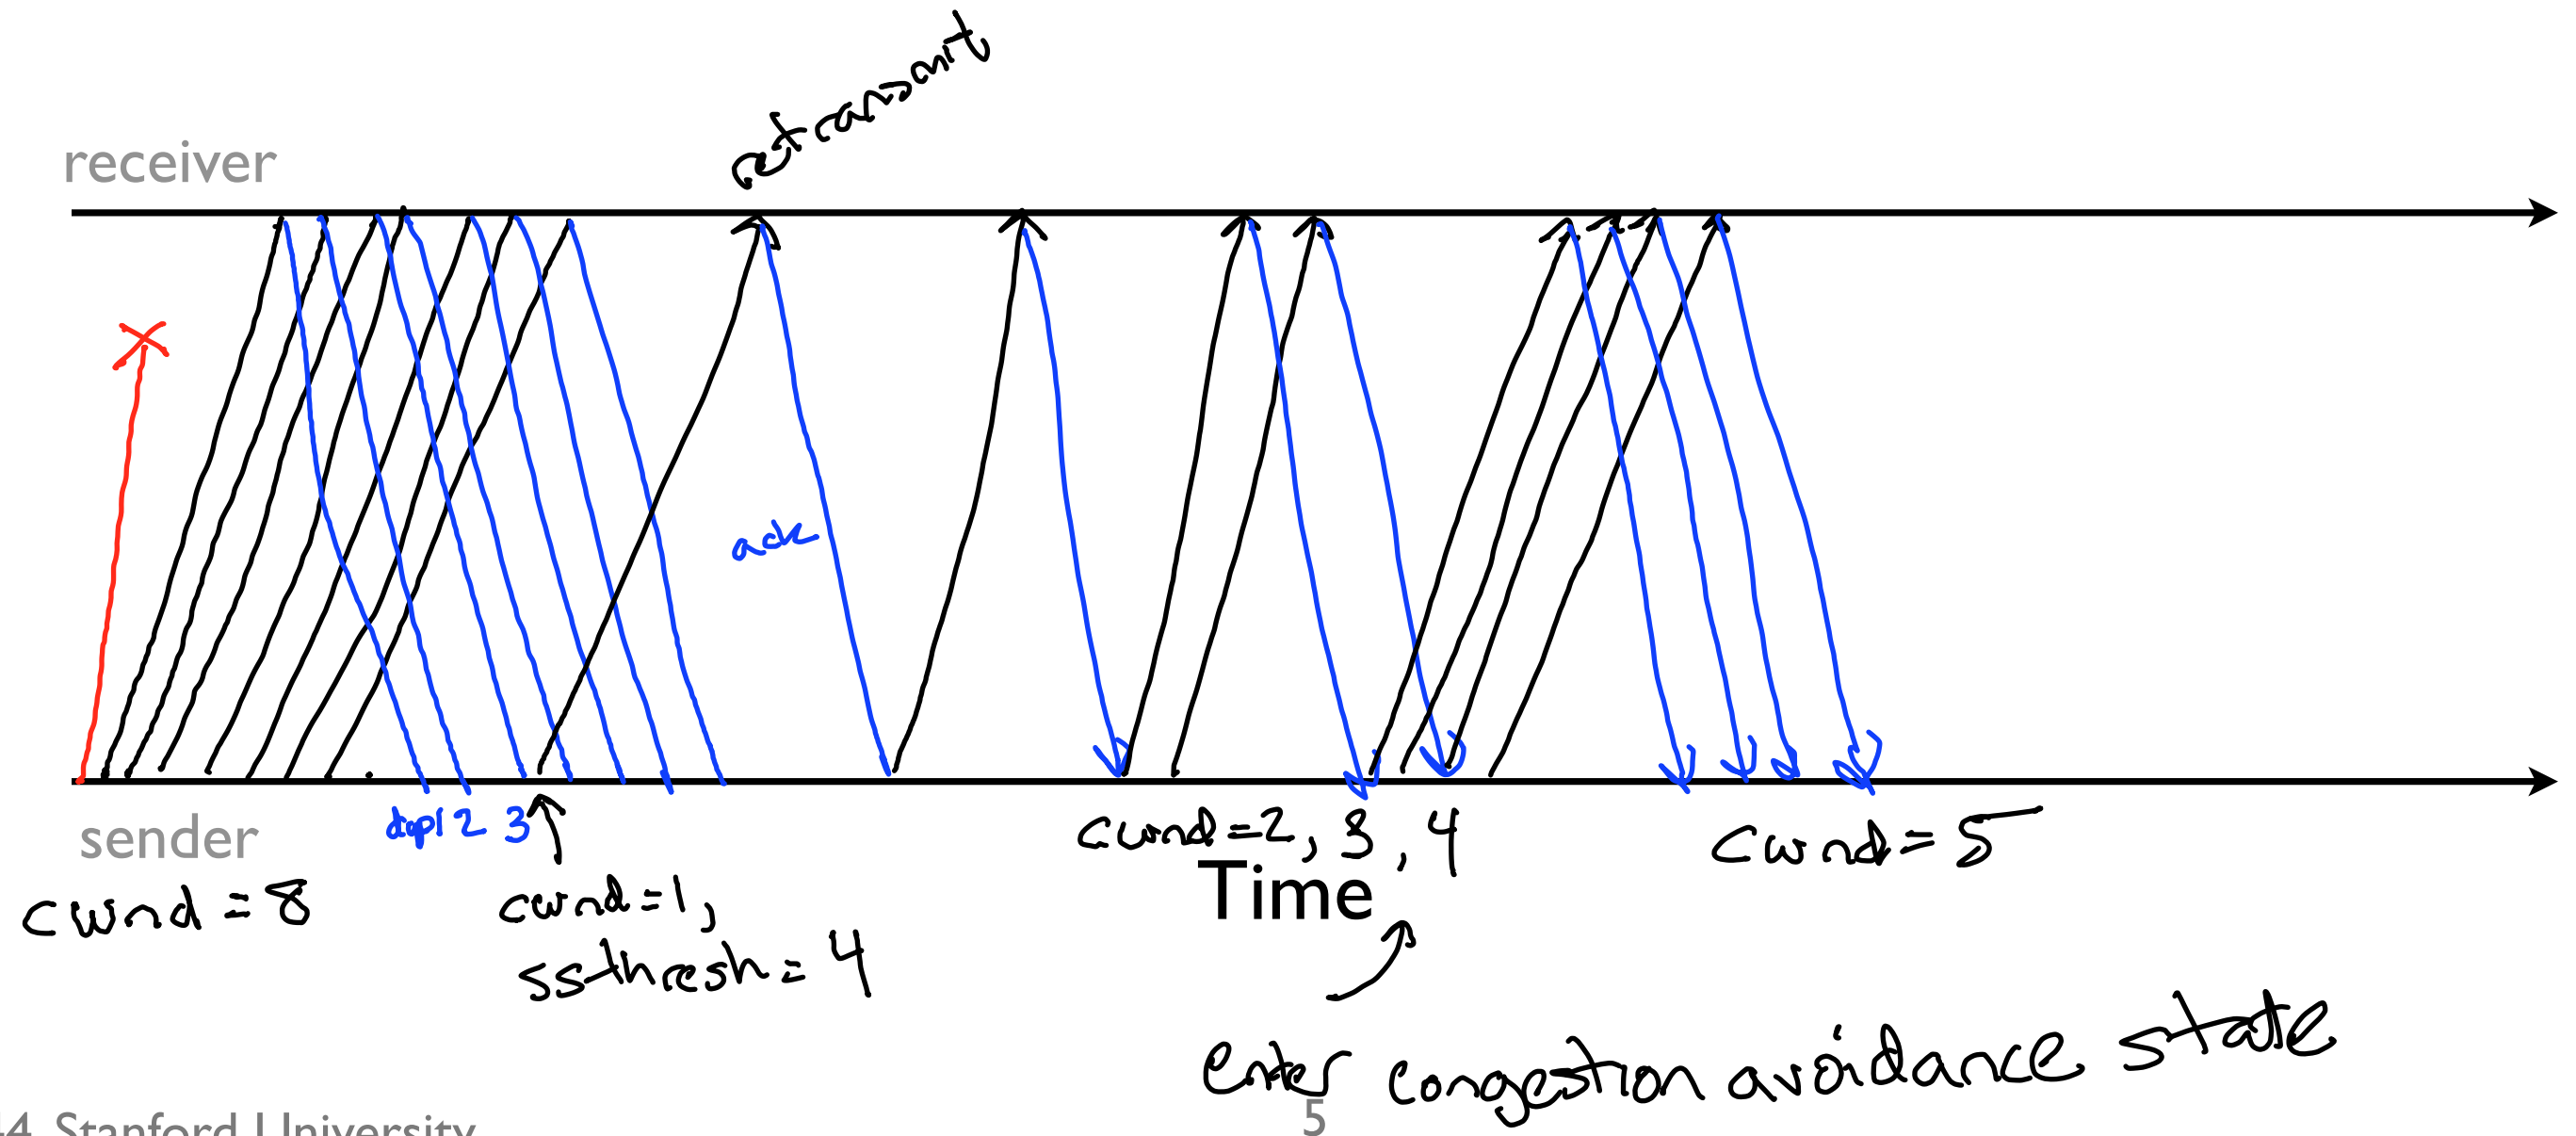
# TCP Tahoe

- On timeout or triple duplicate ack (implies lost packet)
  - ▶ Set threshold to congestion window/2
  - ▶ Set congestion window to 1
  - ▶ Retransmit missing segment (fast retransmit for triple duplicate ack)
  - ▶ Enter slow start state

# TCP Tahoe Behavior



# TCP Tahoe Walkthrough




# TCP Reno

- Same as Tahoe on timeout
- On triple duplicate ack
  - ▶ Set threshold to congestion window/2
  - ▶ Set congestion window to congestion window/2 (fast recovery)
  - ▶ Inflate congestion window size while in fast recovery (fast recovery)
  - ▶ Retransmit missing segment (fast retransmit)
  - ▶ Stay in congestion avoidance state

# TCP Reno

- Same as Tahoe on timeout
- On triple duplicate ack
  - ▶ Set threshold to congestion window/2
  - ▶ Set congestion window to congestion window/2 (fast recovery)
  - ▶ Inflate congestion window size while in fast recovery (fast recovery)
  - ▶ Retransmit missing segment (fast retransmit)
  - ▶ Stay in congestion avoidance state

# TCP Reno






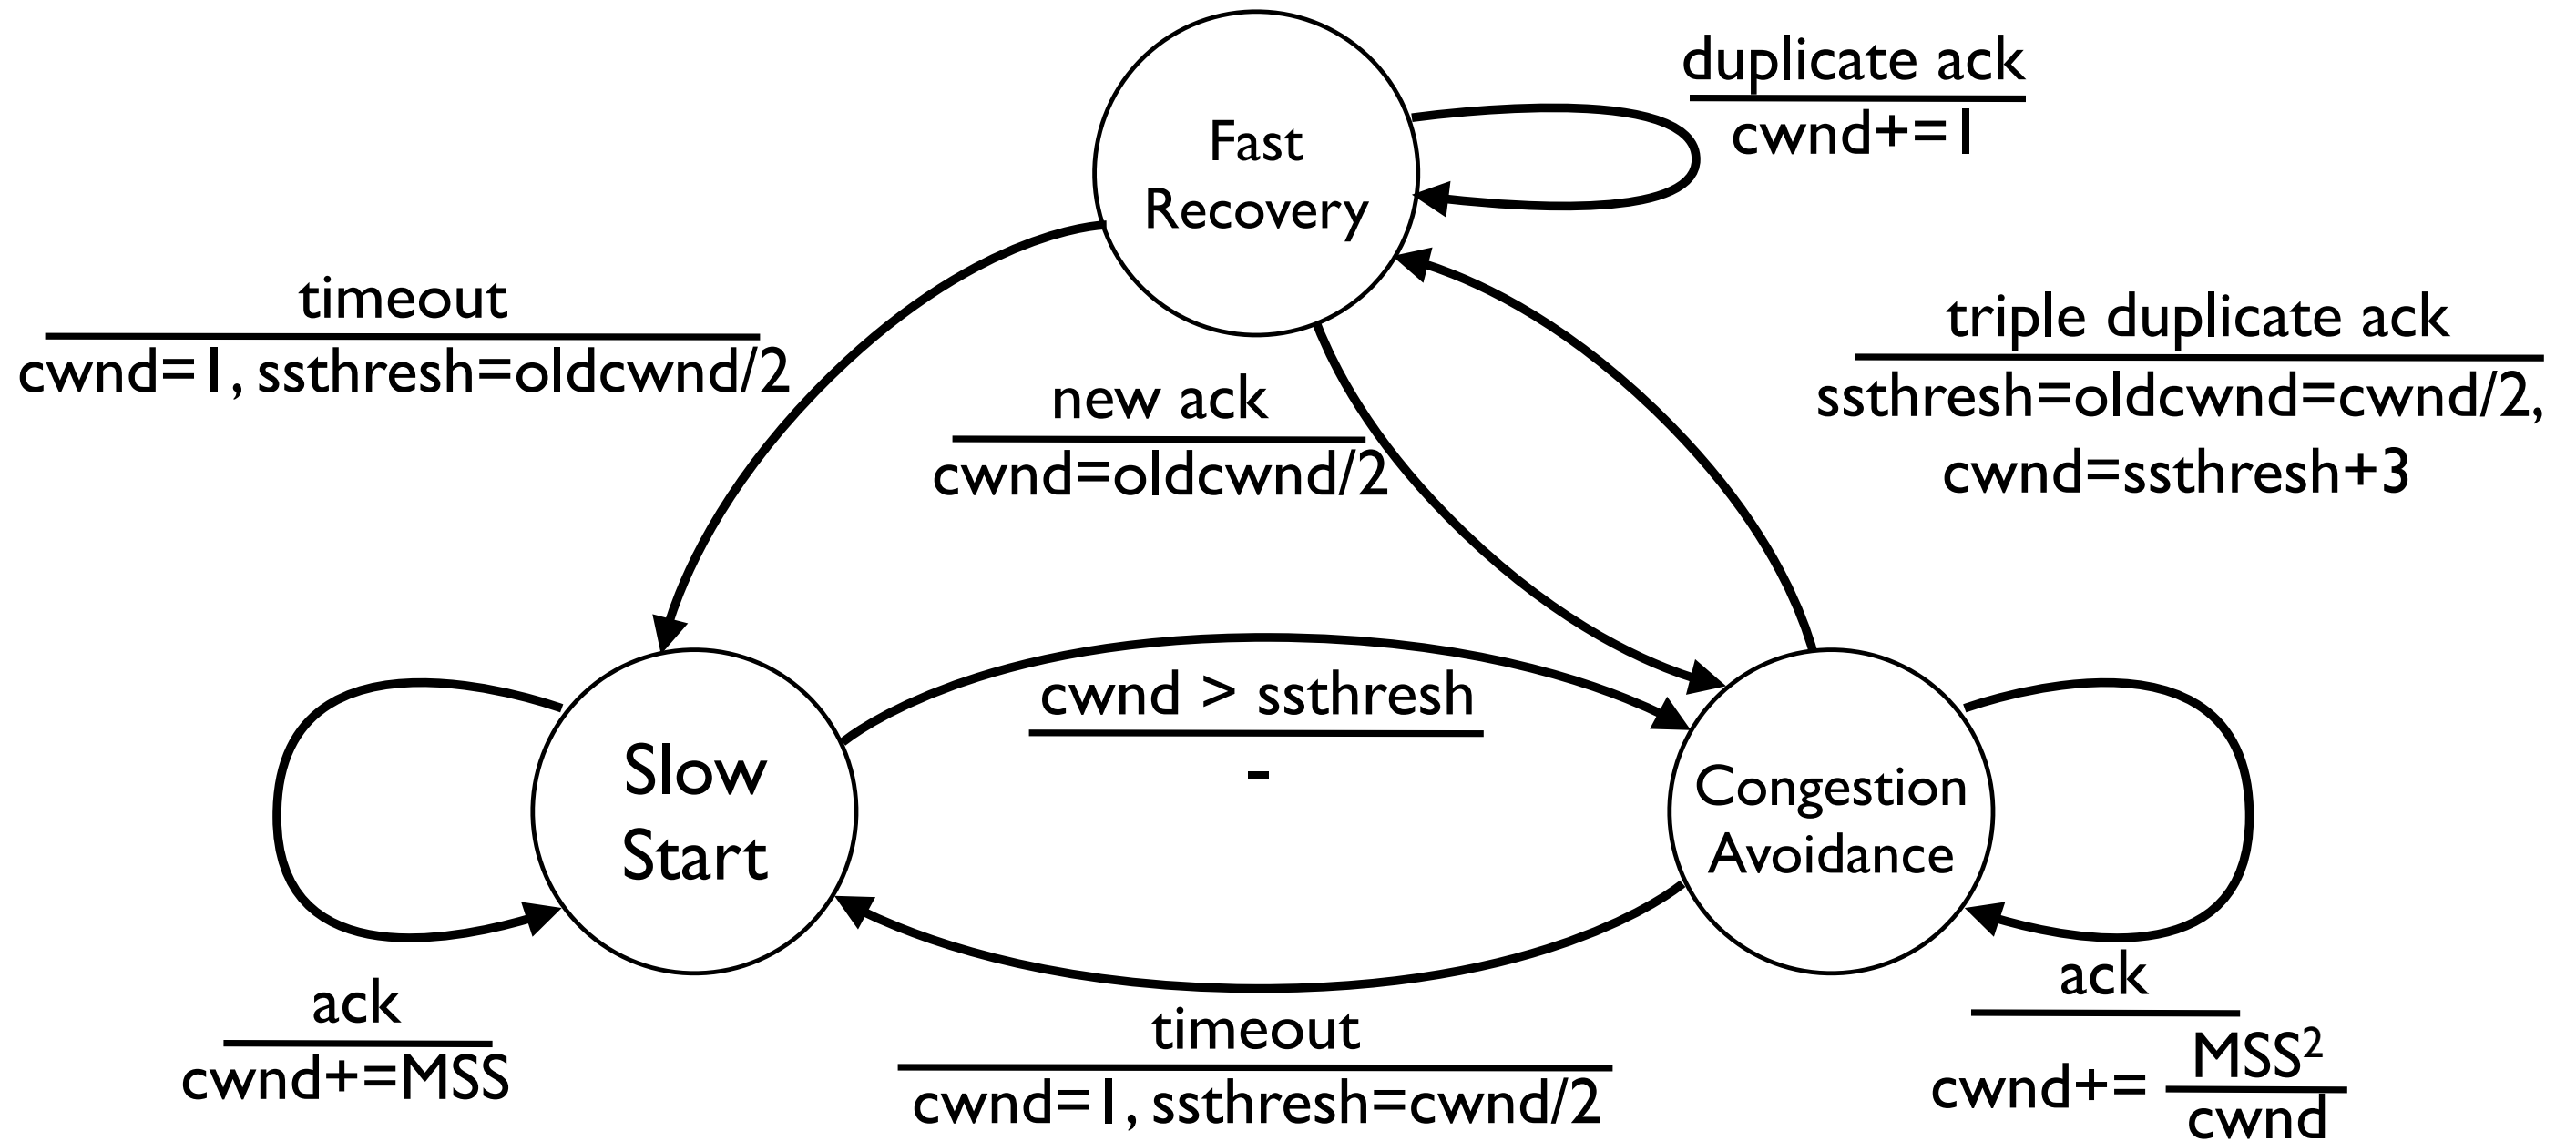
# Congestion Window Inflation

- Problem: it takes a full RTT for a fast retransmitted packet to advance the congestion window
- Could put more packets into network
- Solution: congestion window inflation
  - ▶ While in the fast recovery state (haven't received new acknowledgements), increase congestion window size by 1 for each duplicate acknowledgement, including the initial 3
  - ▶ This happens after halving congestion window size ( $cwnd_{new} = cwnd_{old}/2$ )
  - ▶ End result: after one RTT,  $cwnd_{new}$  is  $3 * cwnd_{old}/2$  but since no new acks yet, this results in sending  $cwnd_{old}/2$  new packets


# TCP Tahoe FSM



# TCP Reno FSM



# TCP Reno Walkthrough



# Congestion Control

- One of the hardest problems in robust networked systems
- Basic approach: additive increase, multiplicative decrease
- Tricks to keep pipe full, improve throughput
  - ▶ Fast retransmit (don't wait for timeout to send lost data)
  - ▶ Congestion window inflation (don't wait an RTT before sending more data)